



CSCI-UA.0480-004

Algorithmic Problem Solving

Lecture 10

Greedy

Bowen Yu
Spring 2018



- HW5 released
 - Bitmask and Trie
 - Due on Sunday
- HW6 released
 - Union find
 - Due on next Wednesday

- Make a greedy choice for a local configuration
- Provable that making local greedy choices leads to a final global optimum.
- A greedy choice is *never* reverted.
 - We do not go back and consider other choices earlier. Otherwise it is not greedy but complete search.
- Greedy algorithm is often ad hoc
- Requires practice + creativity!

There is one row of land with N cells ($N \leq 100$). A cell may contain a crop or is empty. To prevent crops from being eaten by crows, we can install a scarecrow in some cells. A scarecrow installed at cell i protects cell i , cell $i - 1$ and cell $i + 1$ (three consecutive cells). What is the minimum number of scarecrows we need to install to protect all crops?

Sample: # is crop, . is empty

..##..##. => ..**##**..**##**. ans = 2

.#.#.#..# => **.#**..**#**..**#** ans = 3

Solution: For the leftmost crop, place a scarecrow in the cell right to it.

..##..##.	.#.#.#..#
[1]	[1]
..##..##.	.#.#.#..#
[1] [2]	[2]
	.#.#.#..#
	[3]

For any optimal solution, there must be one scarecrow covering the leftmost crop. If it is not right to the crop, we can move it while still being able to get the optimum.

Time: $O(N)$

There is one roll of N cards ($N \leq 100$). The front of a card is white, and the back of a card is black. Initially some cards are facing up while the others are facing down. You can select *exactly* K consecutive cards and flip them together ($1 \leq K \leq N$). Determine if it is possible to flip the cards so that all of them become facing up. If so, also determine the minimum number of flips.

Sample: # is black, . is white

$N = 8, K = 3$

..##.##.

Two moves are needed

..##.##. => ...###.. =>

$N = 8, K = 3$

..##.##.

answer is impossible

Solution: Find the leftmost black card, starting from this card flip K consecutive cards.

$N = 8, K = 3$

..##.##.

1 ..##.##.

....###.

2###.

.....

$N = 8, K = 3$

..##.#.#

1 ..##.#.#

....##.#

2##.#

.....##

3##X

Time: $O(NK)$

This sounds intuitive, but how do we show that:

- If there exists a solution, then this strategy will find it?
- This strategy can give the minimum number of flips?

- Find an intuition for the optimization and come up with a hypothetical greedy strategy
- Examine the strategy on some small cases and look for breakers
- **Prove the strategy**
 - This is the most important step. An unproven greedy strategy is very risky.
 - This is also the most difficult part. Many greedy algorithms are easy to describe, intuitive, but difficult to prove.

Solution: Find the leftmost black card, starting from this card flip K consecutive cards.


How do we show that:

- If there exists a solution, then this strategy will find it?
- This strategy can give the minimum number of flips?

Observation

- There are never more than one flips at a same position in a solution
 - flipping twice at a same position does nothing
- The order of the flips does not matter

Since all flips are at distinct positions, we can order them from left to right and apply them in this order.




The left endpoint of the leftmost flip must be at the first black card. Otherwise we cannot have all cards being white in the end.

Therefore for every optimal solution, the leftmost flip must be the same with the first flip in our greedy strategy. Then we induct for all the flips. ■

There is a room that N people want to book ($N \leq 10^5$). Each person wants to use the room for some time interval $[s, t]$ ($0 \leq s \leq t \leq 10^9$). Find the largest number of people who get can the room without a conflict.

Sample:







$[1, 3]$, $[2, 6]$, $[4, 5]$, $[4, 8]$, $[6, 7]$




Solution:

- Sort the intervals by increasing endTime (tie can be handled arbitrarily).
- Initially latestEndingTime is -1.
- For each interval X, select it if it does not contain latestEndingTime.
- Update latestEndingTime by X.endTime.

Time: $O(N \log N)$



Consider those intervals with a smallest end time t . We must have selected one of them, say X .



- If an optimal solution has one interval ending at t but it is not X , then we can replace it by X .
- If an optimal solution does not have any of those intervals ending at t , then it must contain an interval Y that contains t (otherwise we can add X without conflict). We can thus replace Y by X .

After the replacement, the remaining set of intervals will be the same. Thus our greedy algorithm preserves the opportunity to reach global optimum.

- Prove that any global optimum must contain our local choice
 - e.g. Flips
- Prove that if a global optimum does not contain our local choice, we can make some replacement to let it contain our local choice, while still being able to reach a global optimum
 - e.g. Event Selection


Arrange N non-negative integers onto a binary tree's leaves to minimize

$$\sum_x x \cdot \text{depth}(x)$$

Sample:


[1, 2, 4, 5, 6]

$$\begin{aligned}\text{Cost} &= (1+2)*3 + (4+5+6)*2 \\ &= 39\end{aligned}$$



Huffman Tree Algorithm

- Let S contain all integers. Each integer is a node.
- Select the two smallest integers x, y from S . Create a new node with x, y as its children. Let this new node have value $x + y$. Remove x, y from S . Add $x + y$ to S .
- Repeat until S contains one single integer.



- **Proof of correctness**
 - Show that any optimal tree must have the two smallest integers as a pair of deepest sibling nodes; otherwise we can swap to reduce cost.
 - Show that after processing the two smallest integers, the problem becomes an identical sub-problem where the same strategy applies.
- **Complexity**
 - Using a heap to store S , the total time complexity is $O(N \log N)$

There are N types of coins of distinct values (integers in $[1, 10^9]$). Each type of coin has unlimited supply. Given a total value S ($1 \leq S \leq 10^9$), determine the minimum number of coins needed to compose S . There is always a coin of value 1.

Sample:

coins = [1, 2, 5, 10, 50]

$S = 88$

answer is 7 coins: 50, 10, 10, 10, 5, 2, 1

Incorrect Greedy Strategy: Use as many largest coins as possible

coins = [1, 4, 5, 6]

S = 9

Greedy gets 4 coins: 6, 1, 1, 1

But only 2 are needed: 4, 5

- If we look at real-world coin (bill) values: [1, 2, 5, 10, 50], then the greedy strategy works.
- Someone claims that if we order the coins $C_1 < C_2 < \dots < C_N$ and we have

$$\sum_{1 \leq i < j} C_i \leq C_j$$

then the greedy strategy will work (?)

- e.g. the binary coin system: [1, 2, 4, 8, 16, ...]

Someone claims that if we order the coins $C_1 < C_2 < \dots < C_N$ and we have

$$\sum_{1 \leq i < j} C_i \leq C_j$$

then the greedy strategy will work (?)

- Breaker: coins = [1, 8, 9], S = 24
- Before implementing or proving a “greedy strategy”, try to find small counterexamples that break it.

* Given a set of coins, check if greedy applies is a more complicated topic.

<https://naq17.kattis.com/problems/canonical>

Given a list of N integers ($N \leq 10^5$), determine the length of its longest zigzag subsequence. A sequence is zigzag if $x_1 > x_2 < x_3 > x_4 \dots$

Sample:

$[2, 4, 1, 3, 5] \Rightarrow [\underline{2}, 4, \underline{1}, \underline{3}, 5], \text{ ans} = 3$

$[4, 4, 2, 3, 3, 5, 1] \Rightarrow [4, \underline{4}, \underline{2}, 3, \underline{3}, 5, \underline{1}], \text{ ans} = 4$

Solution: Pick peak and valley alternatively.

peak/valley: last element in an increasing/decreasing streak

Proof: For every increasing streak, we can choose at most two elements (its first and last element). The greedy strategy indeed chooses two elements.

There are N black and N white sticks ($N \leq 100$) of different integer lengths. You are to pair every black stick X with one white stick Y . You are also given two constant integers C (cap), F (fine). If $X.length + Y.length > C$, you have to pay $F * (X.length + Y.length - C)$ dollars. Determine the minimum amount of dollars you have to pay.


Sample:

Black: [1, 3, 4, 5]

White: [2, 4, 4, 5]

$C = 7, F = 10$

answer is 20, we can have



Solution: Pair the shortest white stick with the longest black stick.

Proof:

Let white sticks' lengths be $a_1 \leq a_2 \leq \dots \leq a_n$

Let black sticks' lengths be $b_1 \geq b_2 \geq \dots \geq b_n$


If in an optimal solution, we don't have (a_1, b_1) , then we must have (a_1, b_j) and (a_i, b_1) . In this case we can swap them to obtain a potentially smaller cost. This is because:

$$\begin{array}{l} a_1 \leq a_i \\ | \\ b_1 \geq b_j \end{array}$$

can be
no worse
than

$$\begin{array}{l} a_1 \leq a_i \\ \times \\ b_1 \geq b_j \end{array}$$

$S = a_1 + b_1 + a_i + b_j$ is a constant, consider the line $x + y = S$



We have to pay in the red area.
So we'd better move (x, y) to the middle. Correspondingly, we'd pair a_1 with b_1

There are N jobs ($N \leq 1000$). Job i takes T_i days to complete and has a delay fine F_i . Starting from day 0, we complete the jobs one by one. If a job is started on day d , we have to pay a delay fine of $d * F_i$. Determine the minimum total delay fine we have to pay if we order the jobs optimally.

Sample:

$N = 4$

$T_i = [3, 1, 2, 5]$

$F_i = [4, 1000, 2, 5]$

One optimal order: [2, 1, 3, 4]

Job 2 started on day 0, fine = 0

Job 1 started on day 1, fine = $1 * 4 = 4$

Job 3 started on day 4, fine = $4 * 2 = 8$

Job 4 started on day 6, fine = $6 * 5 = 30$

Total fine = 42

Solution: For a pair of jobs x and y , if $T_x F_y \leq T_y F_x$, then do job x before job y .

Proof:

- Define: $\text{better}(x, y) \Leftrightarrow T_x F_y \leq T_y F_x$
- Trivially for two jobs x and y , the optimal order is $[x, y] \Leftrightarrow \text{better}(x, y)$
- Consider three jobs x, y, z . We want to show $\text{better}(x, y), \text{better}(y, z), \text{better}(x, z) \Rightarrow [x, y, z]$ is optimal

- $\text{better}(x, y), \text{better}(y, z), \text{better}(x, z) \Rightarrow [x, y, z]$ is optimal
 - There are 6 orders: $xyz, xzy, yxz, yzx, zxy, zyx$
 - Since we already know $\text{better}(x, y)$ and $\text{better}(y, z)$, xzy, yxz, zyx cannot be optimal. We only need to discuss xyz, yzx, zxy
 - As an example, we show $\text{cost}(xyz) \leq \text{cost}(yzx)$
 - $\text{cost}(xyz) \leq \text{cost}(zxy)$ is similar.

$$\begin{aligned}
 & \text{cost}(xyz) - \text{cost}(yzx) \\
 &= (T_x F_y + T_x F_z + T_y F_z) - (T_y F_z + T_y F_x + T_z F_x) \\
 &= (T_x F_y - T_y F_x) + (T_x F_z - T_z F_x) \leq 0 \\
 & \quad \text{better}(x, y) \quad \quad \text{better}(x, z)
 \end{aligned}$$

- For three jobs x, y, z , we have proven that $\text{better}(x, y), \text{better}(y, z), \text{better}(x, z) \Rightarrow [x, y, z]$ is optimal. We can extend this to more jobs.
- For three jobs a, b, c with optimal order $[a, b, c]$, and a fourth job d , we can use induction to show $\text{better}(a/b/c, d) \Rightarrow [a, b, c, d]$ is optimal
 - Obviously, $dabc, adbc, abdc$ are not optimal
 - Let “+” denote job combination, $x+y$ is a virtual job (T_x+T_y, F_x+F_y).
 - $\text{better}(a, d)$ and $\text{better}(b, d) \Rightarrow \text{better}(a+b, d) \Rightarrow [a+b, c, d]$ is optimal $\Rightarrow [a, b, c, d]$ is optimal

You are given a 32-bit unsigned integer X and a range $[L, R]$. Find an integer Y ($L \leq Y \leq R$) so that $X \mid Y$ (bitwise OR) is maximum. If there is a tie, choose the smallest Y .

Sample:

$X = 100, L = 50, R = 60 \Rightarrow Y = 59$

X	1100100	R	0111100
Y	0111011	Y	0111011
$X \mid Y$	1111111	L	0110010

Solution:

- Iterate each bit from highest to the lowest
- If the bit is set in X
 - Then we do not need this bit to have a larger OR value
 - But if not setting this bit results in $Y < L$, then we must set this bit
 - We haven't set the lower bits yet. How do we know that Y can be $< L$?
Just try to set all lower bits to 1 and see if $Y < L$.
- If the bit is not set in X
 - Then we need this bit to have a larger OR value
 - But if setting this bit results in $Y > R$, then we must skip this bit

Correctness is implied by the greedy construction and proof is omitted.

You drive a car along a road. Your car initially has a full tank of capacity C . You start at $x=0$ and want to reach $x=L$. There are $N \leq 10^5$ gas stations on the road. The i -th station is at X_i ($0 < X_i < L$) and sells gas for P_i dollars per gallon. At each gas station, you can refill any amount of gas till a full tank. Determine the minimum total cost to reach $x=L$, assuming it costs 1 gallon of fuel to drive 1 unit of distance.

Sample:

$C = 4, L = 6$

$X_1 = 1, P_1 = 3$

$X_2 = 2, P_2 = 2$

$X_3 = 5, P_3 = 1$

Drive to X_2 , fuel 4 \rightarrow 2

At X_2 , buy 1 gallon, fuel 2 \rightarrow 3

Drive to X_3 , fuel 3 \rightarrow 0

At X_3 , buy 1 gallon, fuel 0 \rightarrow 1

Drive to $L=6$, fuel 1 \rightarrow 0