



# CSCI-UA.0480-004 Algorithmic Problem Solving

## Lecture 11

# Divide and Conquer

Bowen Yu  
Spring 2018



- HW6
  - Union find
  - Due tomorrow
- HW7 released
  - Complete search, greedy I
  - Due on Sunday
- HW8 to be released
  - Complete search, greedy II
  - Due on next Wednesday

## Python Power

- A test on modular arithmetic: compute  $a^b \bmod P$ .
- Since  $a$  is large, we can first mod  $a$  by  $P$ .
  - Treat  $a$  as a decimal string.
  - No need to use big integer.
- Then we multiply “ $a \bmod P$ ”  $b$  times with modulus  $P$  to obtain the answer.
- Careful with 32-bit overflow in multiplication.
- Time:  $O(|a| + b)$

## Span Queries

- Use a BBST to maintain the numbers (add, delete).
- Retrieve the leftmost node in the BBST as the min, and the rightmost node in the BBST as the max, in order to compute the span.
- Time:  $O(q \log q)$
- Using the library's heap may have a issue with deletion.
  - `PriorityQueue.remove` is linear!

## Bracket Sequence II

- Similar to homework: If the sequence can be fixed by inserting one bracket, then it must be of the form:
  - $l_1 l_2 l_3 \dots l_k x r_k \dots r_3 r_2 r_1$ 
    - $l_i$  is a left bracket and  $r_i$  is a right bracket, and  $l_i$  matches  $r_i$
    - $x$  is any bracket
- Use a stack to check validness of the sequence. If there is a mismatched bracket, also push it to the stack.
- Finally treat the stack as an array and check if it has the pattern above.

## Students in a Row II

### Solution 1

- Build a doubly linked list for the students. “ml/mr x” is to delete x and then insert at the head/tail of the list.
- Time:  $O(n + q)$

### Solution 2

- Given each student a position value. Initially student x has position x. Maintain the positions in a BBST.
- Record the leftmost and rightmost positions: left=1, right=n.
- For ml, assign “--left” to the student’s position. For mr, assign “++right” to the student’s position.
- For l/r query, find the student in the BBST, and retrieve its predecessor and successor node in the BBST.
- Time:  $O((n + q)\log n)$

## Image Convolution

Store the image and pattern as bitmasks. First let's assume the pattern has no question marks. Suppose each block has  $B$  bits. e.g. for  $B = 8$ :

image	=	00110100		10111011		10101
pattern	=	01101000		01		
		0110100		001		
		011010		0001		

Shift the pattern and compare the bits. We can compare the image against the pattern  $B$  bits at a time. The time complexity is  $O(nmab/B)$ . This  $1/B$  speed boost is required to pass the large subtask.

```
image    = 00110100 | 10 111011 | 10101
pattern = (011010  00)01
```

How to compare when the bitmasks are not aligned? Bit manipulation.

```
image    x = 00110100, y = 10111011
pattern z = 01101000
```

We want the last 6 bits of x and highest 2 bits of y. Use bit manipulation (<<, >>, &) to obtain an integer that is 11010010, then we can compare this against z.



How to handle positions that must be ones? zeroes?

How to handle question marks?

Create a set of pattern bitmasks M1 that has every '#' to be 1.

Create a set of pattern bitmasks M0 that has every '.' to be 1.

Ignore the question marks completely.

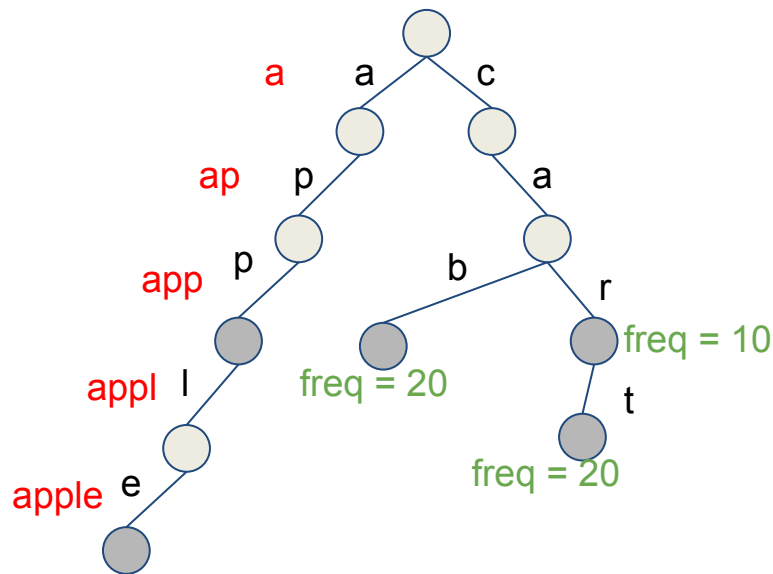
Then a match satisfies:  $M1 \ \& \ image = M1$ ,  $M0 \ \& \ image = 0$

- Using the library's bitset
  - Totally fine.
  - But make sure that you also know how to manipulate integer bitmasks yourselves. We will be doing this a lot on some future topics.
    - e.g. In subset enumeration we typically work with a bitmask with ~20 bits. It is an overkill to use bitset. Using a 32-bit integer can be faster.
- Using the library's big integer.
  - Kind of weird.

## Auto Completion

- Solution to the easy subtask has been directly covered lecture.
- For each Trie node  $X$ , we precompute to which Trie node it auto completes to, say  $\text{ans}(X)$ . Then
  - $\text{ans}(X) =$ 
    - $\text{ans}(Y) \mid Y$  is a child of  $X$ ,  $\text{ans}(Y)$ .string has max frequency and is smallest, or
    - $Y \mid Y$  is a child of  $X$ ,  $Y$ .string has max frequency and is smallest

- For a Trie node X, X.string is an index, not the string itself. Otherwise we exhaust the memory!
  - $O(L^2)$  space, L is 1M, the total number of characters.
- On a tie of max frequency, there is no need to compare the original strings.
  - “cab” has same max frequency as “cart”.
  - The answer is the word following edge “b” because “b” < “r”.



## How to handle multiple tab presses?


- Pressing tab once is essentially jumping from a node  $X$  in the Trie to the node  $X$  completes to.
- Whenever we see a tab, we perform this jump. Since each node  $X$  knows which node it should jump to (precomputed), each jump takes constant time.
- Time
  - $O(|\text{strokes}|)$  per query
  - $O(L)$  in total, where  $L$  is the total number of characters in the input (2M).

## Max XOR on a Tree

- The small subtask gives an array, and is exactly the max XOR subarray problem covered by the lecture.
- Compute  $\text{xor}(i) = a[i] \oplus a[i-1] \oplus a[i-2] \dots a[1]$
- Find the pair  $(i, j)$  where  $\text{xor}(i) \oplus \text{xor}(j-1)$  is maximum using a Trie.

How to extend it to handle a path on the tree?

- Make the tree rooted.
- let  $\text{xor}(x)$  be the XOR of edge weights on the path from  $x$  to the root.
- For a pair of nodes  $x, y$ , the XOR of path  $(x, y)$  is then  $\text{xor}(x) \oplus \text{xor}(y)$ .
- Find the pair  $(x, y)$  using Trie.
- Time:  $O(30n)$ , 30 is the number of bits needed to store  $10^9$ .



You drive a car along a road. Your car initially has a full tank of capacity  $C$ . You start at  $x=0$  and want to reach  $x=L$ . There are  $N \leq 10^5$  gas stations on the road. The  $i$ -th station is at  $X_i$  ( $0 < X_i < L$ ) and sells gas for  $P_i$  dollars per gallon. At each gas station, you can refill any amount of gas till a full tank. Determine the minimum total cost to reach  $x=L$ , assuming it costs 1 gallon of fuel to drive 1 unit of distance.

### Sample:

$C = 4, L = 6$

$X_1 = 1, P_1 = 3$

$X_2 = 2, P_2 = 2$

$X_3 = 5, P_3 = 1$

Drive to  $X_2$ , fuel 4  $\rightarrow$  2

At  $X_2$ , buy 1 gallon, fuel 2  $\rightarrow$  3

Drive to  $X_3$ , fuel 3  $\rightarrow$  0

At  $X_3$ , buy 1 gallon, fuel 0  $\rightarrow$  1

Drive to  $L=6$ , fuel 1  $\rightarrow$  0



## Solution

- If we can find a station that has cheaper gas and within the range of a full tank's distance
  - If with the remaining fuel we can drive to it, then drive to it
  - Otherwise, buy just an amount of fuel so that we can reach that station with zero fuel left
- If we cannot find such a cheaper station, then refill to a full tank at the current station. Then drive to the next station.

Proof is implied by the greedy decision.

Straightforward implementation takes  $O(N^2)$ . Can be optimized to  $O(N)$ .

## Divide and Conquer (D&C)

- Divide the problem into smaller subproblems of a same structure and solve each subproblem
- Merge the answers to the subproblems to form the answer to the original problem.
- Many data structures use D&C.

## • Merge Sort

- Split the array into two equal halves. Merge sort each half. Then merge the sorted half arrays.
- Merging two sorted arrays takes linear time. We repeatedly take the smaller of the two arrays' heads.
- $T(n) = 2T(n/2) + O(n) = O(n \log n)$

## • Quick Sort

- Choose a pivot element. Place the elements smaller than the pivot on the left, and the elements larger than (or equal to) the pivot element on the right. This rearrangement takes  $O(n)$ .
- Quick sort the left and right groups individually.
- Ideally, the pivot splits the array into two equal halves.  $T(n) = 2T(n/2) + O(n) = O(n \log n)$ 
  - But if we choose a bad pivot, this can downgrades to  $O(n^2)$

## Master Theorem

### *Theorem 4.1 (Master theorem)*

Let  $a \geq 1$  and  $b > 1$  be constants, let  $f(n)$  be a function, and let  $T(n)$  be defined on the nonnegative integers by the recurrence

$$T(n) = aT(n/b) + f(n) ,$$

where we interpret  $n/b$  to mean either  $\lfloor n/b \rfloor$  or  $\lceil n/b \rceil$ . Then  $T(n)$  has the following asymptotic bounds:

1. If  $f(n) = O(n^{\log_b a - \epsilon})$  for some constant  $\epsilon > 0$ , then  $T(n) = \Theta(n^{\log_b a})$ .
2. If  $f(n) = \Theta(n^{\log_b a})$ , then  $T(n) = \Theta(n^{\log_b a} \lg n)$ .
3. If  $f(n) = \Omega(n^{\log_b a + \epsilon})$  for some constant  $\epsilon > 0$ , and if  $af(n/b) \leq cf(n)$  for some constant  $c < 1$  and all sufficiently large  $n$ , then  $T(n) = \Theta(f(n))$ . ■

## Master Theorem (Intuitively)

$$T(n) = aT(n/b) + f(n)$$

- We want to compare  $O(n^{\log_b a})$  and  $f(n)$
- If one grows faster than the other, then the total complexity is that function.
- If both grows equally fast, then it's  $O(n^{\log_b a} \log n)$

$$T(n) = 2T(n/2) + O(1) \Rightarrow O(n)$$

$$T(n) = 2T(n/2) + O(n^2) \Rightarrow O(n^2)$$

$$T(n) = 3T(n/2) + O(n) \Rightarrow O(n^{\log_2 3})$$

$$T(n) = 4T(n/2) + O(n^2) \Rightarrow O(n^2 \log n)$$

$$T(n) = T(2n/3) + O(1) \Rightarrow O(\log n)$$

We want to compute  $(a^b \bmod P)$ , where  $P = 10^9 + 7$ ,  $0 \leq a, b \leq 10^9$ .

- If  $b$  is small (e.g.  $\leq 10^6$ ), we can just multiply  $b$  times. But for  $b=10^9$ ,  $O(b)$  will TLE.

## Solution

- We can compute  $a^1, a^2, a^4, a^8$  and so on in  $\log(b)$  time.
- Then we can take the powers matching  $b$ 's binary representation.
  - e.g.  $b = 5 = (101)_2$ , then  $a^b = a^1 * a^4$

```

1 int modPower(int a, int b, int P) {
2     int ans = 1;
3     while (b) {
4         if (b & 1) ans = (long long)ans * a % P;
5         b >>= 1;
6         a = (long long)a * a % P;
7     }
8     return ans;
9 }

```

- The while loop is checking each bit of b, from the least significant to the most significant.
- After checking each bit, we square a, so that a becomes  $a^2$ ,  $a^4$ ,  $a^8$ , and so on.



We want to use 32-bit integers without cast to 64-bit to compute  $(a * b \bmod P)$ , where  $P = 10^9 + 7$ ,  $0 \leq a, b \leq 10^9$ .

```
1 int modMultiply(int a, int b, int P){
2     int ans = 0;
3     while (b) {
4         if (b & 1) ans = (ans + a) % P;
5         b >>= 1;
6         a = a * 2 % P;
7     }
8     return ans;
9 }
```

- We are adding  $a$ ,  $2a$ ,  $4a$ ,  $8a$ , and so on if they exist in the binary representation of  $b$ .

We have Fibonacci numbers  $\text{fib}(0) = 0$ ,  $\text{fib}(1) = 1$ , and  $\text{fib}(i) = \text{fib}(i-1) + \text{fib}(i-2)$  for  $i \geq 2$ .

Compute  $\text{fib}(n) \bmod P = 10^9+7$ ,  $n \leq 10^9$


- Again, we cannot afford  $O(n=10^9)$  to linearly compute it.

$$M = \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix}$$

$$M \cdot \begin{bmatrix} fib(i-1) \\ fib(i) \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} fib(i-1) \\ fib(i) \end{bmatrix} = \begin{bmatrix} fib(i) \\ fib(i+1) \end{bmatrix}$$


$$M^{n-1} \cdot \begin{bmatrix} fib(0) \\ fib(1) \end{bmatrix} = \begin{bmatrix} fib(n-1) \\ fib(n) \end{bmatrix}$$

We can compute  $M^{n-1}$  using logarithm power (with modulus). The time complexity is  $O(\log n * k^3)$ , where the matrix is of size  $k$ -by- $k$ .



Count how many ways to use two types (L shape and single cell) of blocks to construct a wall of size  $2 \times N$  ( $N \leq 10^9$ ). L shape can be rotated.


Sample: 11 ways for  $N = 3$



Let  $f(n)$  be the number of ways to build a wall of  $2 \times n$ .

We have  $f(n) = f(n-1) + 4f(n-2) + 2f(n-3)$ , for  $n \geq 3$ .

Boundary:  $f(0) = 1$ ,  $f(1) = 1$ ,  $f(2) = 5$



$$f(n) = f(n-1) + 4f(n-2) + 2f(n-3)$$

$$M = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 2 & 4 & 1 \end{bmatrix} \quad M \cdot \begin{bmatrix} f(n-3) \\ f(n-2) \\ f(n-1) \end{bmatrix} = \begin{bmatrix} f(n-2) \\ f(n-1) \\ f(n) \end{bmatrix}$$

Time complexity:  $O(\log n * 3^3)$

- Form a monotonic search space.
- Check the middle point of the search space.
- Reduce the search space by half.

Given an array of  $N$  integers ( $N \leq 10^5$ ), and  $Q$  queries. Each query has a number  $x$  and asks if  $x$  is in the array.

- $Q = 1$ : Linear scan
- $Q = 10^5$ : Sort the array before all the queries, and binary search for each query. Time:  $O(Q \log N)$ 
  - Search space is index range  $l = 0, r = N-1$
  - Pick the middle index  $m = (l + r)/2$ , if  $arr[m] < x$ , then search space becomes  $[m + 1, r]$ ; Otherwise, search space becomes  $[l, m - 1]$ .




- There are binary search util functions for binary searching on a sorted array.
- However we must know how to write a binary search ourselves because for many problems the search space is not just an array.

```

1 int binarySearch(int l, int r) {
2     while (l <= r) {
3         int m = (l + r) / 2;
4         if (condition(m)) r = m - 1;
5         else r = m + 1;
6     }
7     return l; // or r
8 }

```



Want the smallest  $i$  so that  $\text{condition}(i) = \text{true}$ , pick  $l$ .

Want the largest  $i$  so that  $\text{condition}(i) = \text{false}$ , pick  $r$ .

- Be careful with binary search's search space reduction and exit condition.
  - Some people use `while (l < r)` in place of `while (l <= r)`
  - Some people use `l = m`, `r = m` when reducing the search space by half.
- The while loop may not terminate if this is not set correctly!
- If `l` and `r` are floating point numbers, then we can just loop for a constant number of times (e.g. 100 times).
  - Each time the range is cut into half, and after  $k$  times we have a precision of  $2^{-k}$

- Sometimes for a problem that asks for an answer  $X$ , directly computing  $X$  can be difficult.
- However, given  $X$ , checking if  $X$  is an answer is easy.
- If the problem's solution space is monotonic, then we can binary search the solution space.
  - Monotonic solution space example: If  $X$  works, then every  $Y > X$  also works.

Given a count  $k$  ( $k \leq 10^9$ ), find the smallest integer  $n$ , so that  $n!$  has at least  $k$  trailing zeroes.

- Directly finding such a smallest  $n$  is difficult (or tedious if not impossible).
- However, given  $n$  calculating how many trailing zeroes  $n!$  has is straightforward.

- The number of trailing zeroes in  $n!$  only depends on how many factors of 2's and 5's  $n!$  has.
- For  $n!$  the number of 2's is always greater than the number of 5's, so we can simply count the number of 5's.

	5	10	15	20	25	30	...	50	...	125
$5^1$ (1st 5)	✓	✓	✓	✓	✓	✓		✓		✓
$5^2$ (2nd 5)					✓			✓		✓
$5^3$ (3rd 5)										✓

Therefore trailing zeroes in  $n!$  equals  $\text{floor}(n/5^1) + \text{floor}(n/5^2) + \text{floor}(n/5^3) + \dots$

## Solution

- Binary search on the answer  $n$
- Initial range is  $[l, r] = [1, 10^9]$
- Let  $m = (l + r)/2$ , count the number of trailing zeroes in  $m!$ , say  $t$ .
  - If  $t \geq k$ :  $[l, r] \leftarrow [l, m - 1]$
  - else:  $[l, r] \leftarrow [m + 1, r]$
- Here, condition is  $t \geq k$ , we want the *smallest* value where the condition holds. So finally we return the last value of  $l$ .

Time:  $O(\log^2 n)$ . We binary search on  $n$  which is  $O(\log n)$ . Counting the number of trailing zeroes for a given  $m$  is also  $O(\log n)$ .